



Mid-term programming assignments

Parallel Programming

What to do

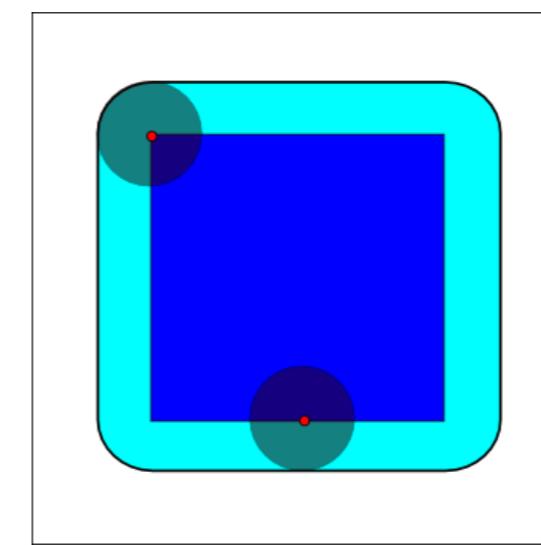
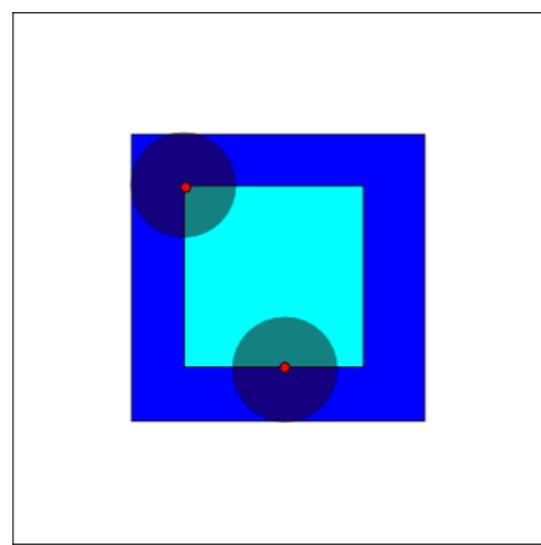
- It's possible to implement the assignments in pairs
- implement a programming assignment using one of the approaches presented (parallelization/vectorization). Implement a sequential and a parallel version.
- Write a tech report with performance analysis (speedup) and prepare a presentation
- Keep everything on a public Github/Bitbucket/Gitlab

Image reader

- Multithread JPEG image reader
 - read X images from a directory
 - Keep the uncompressed bitmaps (or the compressed stream) in an appropriate data structure, that manages access to the images
 - OpenMP does not allow to implement a non-blocking version since it requires asynchronous multi-threading (ie.. CUDA allows it).

Morphological image processing

- Image processing
 - CUDA is a good candidate
 - See https://en.wikipedia.org/wiki/Mathematical_morphology



Bigrams / trigrams

- Compute histograms of bigrams/trigrams on texts (eg. Wikipedia / Gutenberg project documents)
 - Consider bigrams/trigrams of characters and words
 - If needed re-use more and more times the collected data to create a large enough corpus



Password decryption

- Decrypt passwords encrypted using crypt (DES)
 - Consider 8 characters lengths in the set [a-zA-Z0-9.]
 - CUDA does not provide the crypt C-library function.
Must use a special implementation (check Moodle)
 - It's a (slowed down) search problem.
 - Can choose to attack a specific password in a list (in different positions), or decrypt a full list of passwords (in this case reduce the search space)
- Can consider only dates or common 8-chars passwords

Pattern recognition



- Search a given time series within a larger and longer set of time series
 - Get time series from existing machine learning datasets
 - Use SAD as the metric to evaluate the match



α_1	α_2	α_3	α_4	α_5	α_6	α_7	α_8	α_9	α_{10}	α_{11}	α_{12}	α_{13}	α_{14}	α_{15}	α_{16}
β_1	β_2	β_3	β_4	β_5	β_6	β_7	β_8	β_9	β_{10}	β_{11}	β_{12}	β_{13}	β_{14}	β_{15}	β_{16}



γ_1	γ_2	γ_3	γ_4	γ_5	γ_6
δ_1	δ_2	δ_3	δ_4	δ_5	δ_6

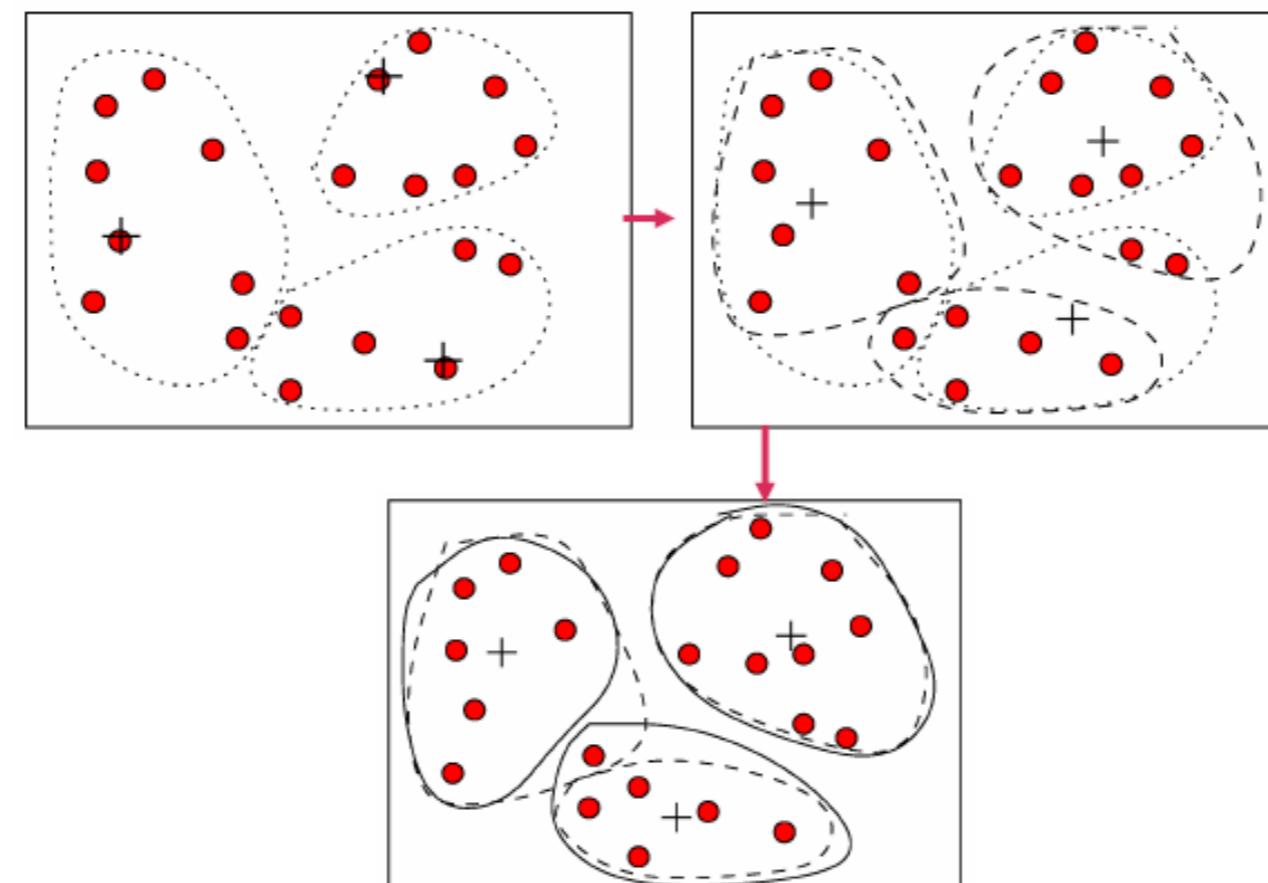
$$|\alpha_1 - \gamma_1| + \dots + |\alpha_6 - \gamma_6| + \\ |\beta_1 - \delta_1| + \dots + |\beta_6 - \delta_6|$$

$$|\alpha_{11} - \gamma_1| + \dots + |\alpha_{16} - \gamma_6| + \\ |\beta_{11} - \delta_1| + \dots + |\beta_{16} - \delta_6|$$



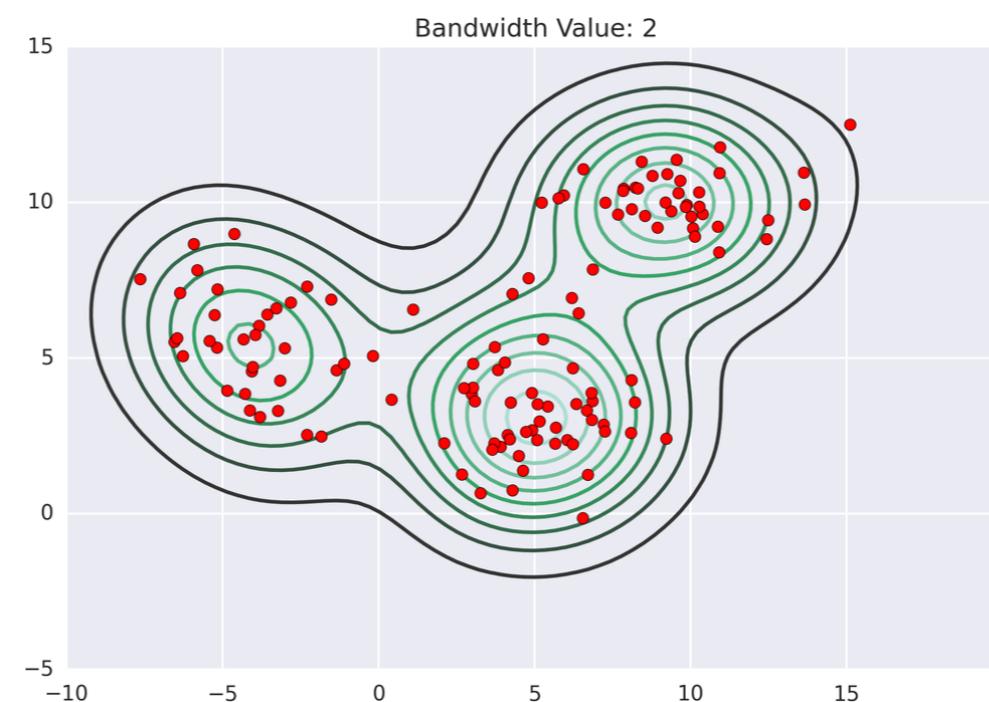
k-means

- Implement a parallel version of k-means clustering
 - There's no need to implement K-means++ centroid assignments
 - Feel free to chose to operate on 2D or larger dimensionality vectors
 - **Esempio**



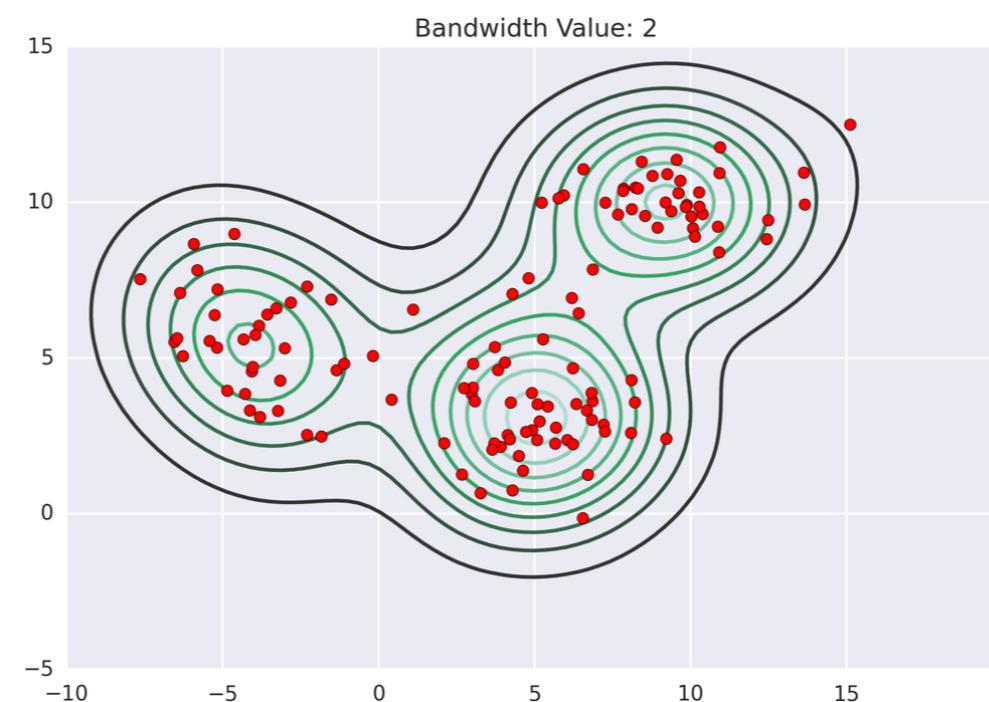
Mean shift clustering

- Implement a parallel version of mean-shift clustering. This clustering doesn't require to specify the number of desired clusters (it uses KDE, so needs a bandwidth parameter) but its complexity is $O(N^2)$
- Can be used to segment images



Mean shift clustering

- Implement a parallel version of mean-shift clustering. This clustering doesn't require to specify the number of desired clusters (it uses KDE, so needs a bandwidth parameter) but its complexity is $O(N^2)$
- Can be used to segment images



ANN retrieval

- Consider high dimensionality vectors (eg. 128)
- Given a specified query find the *k nearest neighbors* in a database of vectors
 - Can use LSH for approximate indices, eg. using mlpack or LSHkit
 - Can get sample datasets from <http://corpus-texmex.irisa.fr>

Renderer



- Implement a renderer that draws (semi) transparent circles (or other shapes). Circles have 3D coordinates and the order along Z axis matters for the correct rendering
 - A pixel belongs to a circle if its center is within the circle.

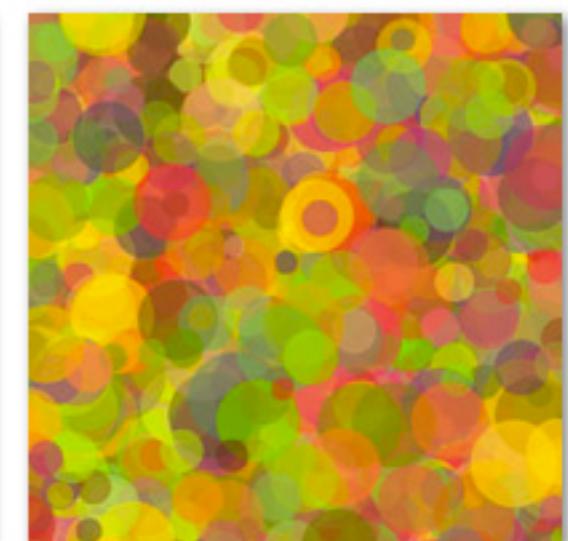
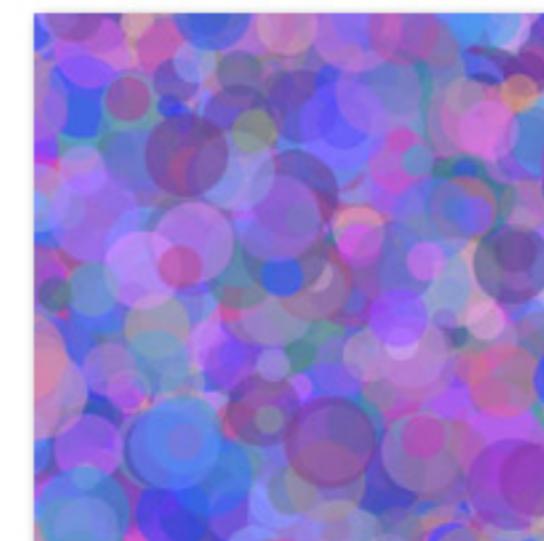
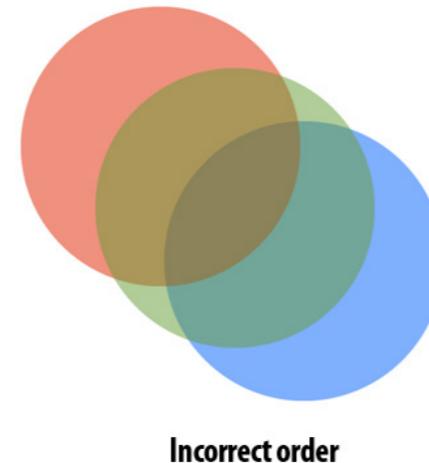
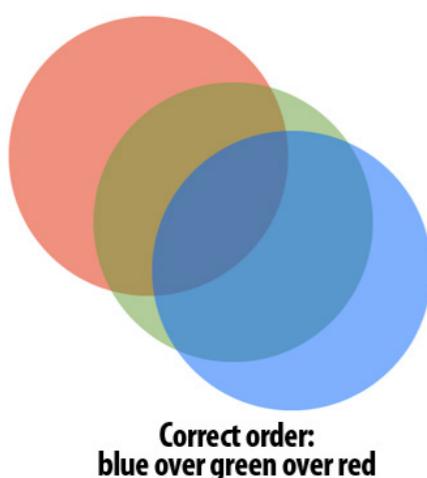
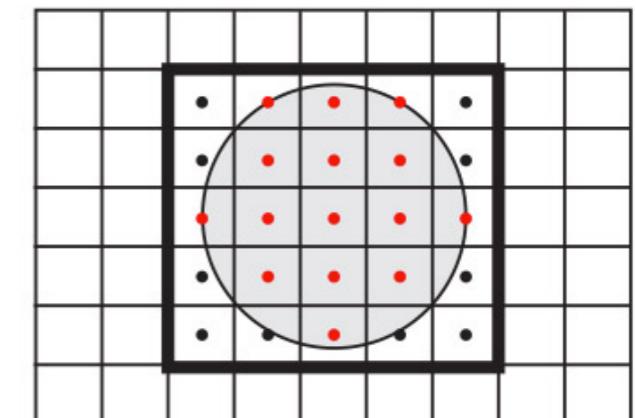




Image composition

- Implement a parallel alpha composition of images
- Useful to create image augmentation to train Convolutional Neural Networks, e.g. for detection (see Blend augmentation in IMgAug library)
- Can parallelize applying the same object to different images, or in different positions in the same image



Random maze solver

- Find the exit from a maze using the random movement of a particle.
 - Start a large number of particles, move them randomly bouncing on the walls
 - Backtrack the first particle to get out of the maze to find the exit

